

Lab08 Mini Project

Elsa Quillin (PID: A69043008)

Table of contents

Background	1
Data import	1
Principal Component Analysis	5
Hierarchical Clustering	13
Combining methods	15
Prediction	18

Background

The goal of this mini-project is for you to explore a complete analysis using the unsupervised learning techniques covered in the last class. We will extend what you've learned by combining PCA as a preprocessing step to clustering using data that consist of measurements of cell nuclei of human breast masses. This expands on our RNA-Seq analysis from last day.

The data itself comes from the Wisconsin Breast Cancer Diagnostic Data Set first reported by K. P. Benne and O. L. Mangasarian: "Robust Linear Programming Discrimination of Two Linearly Inseparable Sets".

Values in this data set describe characteristics of the cell nuclei present in digitized images of a fine needle aspiration (FNA) of a breast mass.

Data import

```
fna.data <- "https://bioboot.github.io/bimm143_S20/class-material/WisconsinCancer.csv"  
wisc.df <- read.csv(fna.data, row.names = 1)  
head(wisc.df)
```

		diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	
842302	M	17.99	10.38	122.80	1001.0		
842517	M	20.57	17.77	132.90	1326.0		
84300903	M	19.69	21.25	130.00	1203.0		
84348301	M	11.42	20.38	77.58	386.1		
84358402	M	20.29	14.34	135.10	1297.0		
843786	M	12.45	15.70	82.57	477.1		
		smoothness_mean	compactness_mean	concavity_mean	concave.points_mean		
842302		0.11840	0.27760	0.3001	0.14710		
842517		0.08474	0.07864	0.0869	0.07017		
84300903		0.10960	0.15990	0.1974	0.12790		
84348301		0.14250	0.28390	0.2414	0.10520		
84358402		0.10030	0.13280	0.1980	0.10430		
843786		0.12780	0.17000	0.1578	0.08089		
		symmetry_mean	fractal_dimension_mean	radius_se	texture_se	perimeter_se	
842302		0.2419		0.07871	1.0950	0.9053	8.589
842517		0.1812		0.05667	0.5435	0.7339	3.398
84300903		0.2069		0.05999	0.7456	0.7869	4.585
84348301		0.2597		0.09744	0.4956	1.1560	3.445
84358402		0.1809		0.05883	0.7572	0.7813	5.438
843786		0.2087		0.07613	0.3345	0.8902	2.217
		area_se	smoothness_se	compactness_se	concavity_se	concave.points_se	
842302		153.40	0.006399	0.04904	0.05373	0.01587	
842517		74.08	0.005225	0.01308	0.01860	0.01340	
84300903		94.03	0.006150	0.04006	0.03832	0.02058	
84348301		27.23	0.009110	0.07458	0.05661	0.01867	
84358402		94.44	0.011490	0.02461	0.05688	0.01885	
843786		27.19	0.007510	0.03345	0.03672	0.01137	
		symmetry_se	fractal_dimension_se	radius_worst	texture_worst		
842302		0.03003		0.006193	25.38	17.33	
842517		0.01389		0.003532	24.99	23.41	
84300903		0.02250		0.004571	23.57	25.53	
84348301		0.05963		0.009208	14.91	26.50	
84358402		0.01756		0.005115	22.54	16.67	
843786		0.02165		0.005082	15.47	23.75	
		perimeter_worst	area_worst	smoothness_worst	compactness_worst		
842302		184.60	2019.0	0.1622	0.6656		
842517		158.80	1956.0	0.1238	0.1866		
84300903		152.50	1709.0	0.1444	0.4245		
84348301		98.87	567.7	0.2098	0.8663		
84358402		152.20	1575.0	0.1374	0.2050		
843786		103.40	741.6	0.1791	0.5249		
		concavity_worst	concave.points_worst	symmetry_worst			

842302	0.7119	0.2654	0.4601
842517	0.2416	0.1860	0.2750
84300903	0.4504	0.2430	0.3613
84348301	0.6869	0.2575	0.6638
84358402	0.4000	0.1625	0.2364
843786	0.5355	0.1741	0.3985
	fractal_dimension_worst		
842302		0.11890	
842517		0.08902	
84300903		0.08758	
84348301		0.17300	
84358402		0.07678	
843786		0.12440	

```
wisc.data <- wisc.df[,-1]
head(wisc.data)
```

	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean
842302	17.99	10.38	122.80	1001.0	0.11840
842517	20.57	17.77	132.90	1326.0	0.08474
84300903	19.69	21.25	130.00	1203.0	0.10960
84348301	11.42	20.38	77.58	386.1	0.14250
84358402	20.29	14.34	135.10	1297.0	0.10030
843786	12.45	15.70	82.57	477.1	0.12780
	compactness_mean	concavity_mean	concave.points_mean	symmetry_mean	
842302	0.27760	0.3001		0.14710	0.2419
842517	0.07864	0.0869		0.07017	0.1812
84300903	0.15990	0.1974		0.12790	0.2069
84348301	0.28390	0.2414		0.10520	0.2597
84358402	0.13280	0.1980		0.10430	0.1809
843786	0.17000	0.1578		0.08089	0.2087
	fractal_dimension_mean	radius_se	texture_se	perimeter_se	area_se
842302		0.07871	1.0950	0.9053	8.589 153.40
842517		0.05667	0.5435	0.7339	3.398 74.08
84300903		0.05999	0.7456	0.7869	4.585 94.03
84348301		0.09744	0.4956	1.1560	3.445 27.23
84358402		0.05883	0.7572	0.7813	5.438 94.44
843786		0.07613	0.3345	0.8902	2.217 27.19
	smoothness_se	compactness_se	concavity_se	concave.points_se	
842302	0.006399	0.04904	0.05373		0.01587
842517	0.005225	0.01308	0.01860		0.01340
84300903	0.006150	0.04006	0.03832		0.02058

```

84348301      0.009110      0.07458      0.05661      0.01867
84358402      0.011490      0.02461      0.05688      0.01885
843786       0.007510      0.03345      0.03672      0.01137
          symmetry_se fractal_dimension_se radius_worst texture_worst
842302        0.03003       0.006193      25.38       17.33
842517        0.01389       0.003532      24.99       23.41
84300903      0.02250       0.004571      23.57       25.53
84348301      0.05963       0.009208      14.91       26.50
84358402      0.01756       0.005115      22.54       16.67
843786       0.02165       0.005082      15.47       23.75
          perimeter_worst area_worst smoothness_worst compactness_worst
842302         184.60      2019.0       0.1622      0.6656
842517         158.80      1956.0       0.1238      0.1866
84300903      152.50      1709.0       0.1444      0.4245
84348301        98.87       567.7       0.2098      0.8663
84358402      152.20      1575.0       0.1374      0.2050
843786         103.40       741.6       0.1791      0.5249
          concavity_worst concave.points_worst symmetry_worst
842302         0.7119       0.2654       0.4601
842517         0.2416       0.1860       0.2750
84300903      0.4504       0.2430       0.3613
84348301        0.6869       0.2575       0.6638
84358402        0.4000       0.1625       0.2364
843786         0.5355       0.1741       0.3985
          fractal_dimension_worst
842302           0.11890
842517           0.08902
84300903          0.08758
84348301           0.17300
84358402           0.07678
843786           0.12440

```

```
diagnosis <- as.factor(wisc.df$diagnosis)
```

Q1. How many observations are in this dataset?

```
dim(wisc.data)
```

```
[1] 569 30
```

There are 30 observations for each sample

Q2. How many of the observations have a malignant diagnosis?

```
sum(diagnosis == "M")
```

```
[1] 212
```

212 of the observations have a malignant diagnosis, so 357 are benign

Q3. How many variables/features in the data are suffixed with `_mean`?

```
length(grep("_mean", colnames(wisc.data)))
```

```
[1] 10
```

There are 10 columns that are suffixed with “`_mean`”

Principal Component Analysis

The main function in base R for PCA is called `prcomp()`. An optional argument `scale()` should nearly always be switched to `scale=TRUE` for this function.

```
colMeans(wisc.data)
```

radius_mean	texture_mean	perimeter_mean
1.412729e+01	1.928965e+01	9.196903e+01
area_mean	smoothness_mean	compactness_mean
6.548891e+02	9.636028e-02	1.043410e-01
concavity_mean	concave.points_mean	symmetry_mean
8.879932e-02	4.891915e-02	1.811619e-01
fractal_dimension_mean	radius_se	texture_se
6.279761e-02	4.051721e-01	1.216853e+00
perimeter_se	area_se	smoothness_se
2.866059e+00	4.033708e+01	7.040979e-03
compactness_se	concavity_se	concave.points_se
2.547814e-02	3.189372e-02	1.179614e-02
symmetry_se	fractal_dimension_se	radius_worst
2.054230e-02	3.794904e-03	1.626919e+01
texture_worst	perimeter_worst	area_worst
2.567722e+01	1.072612e+02	8.805831e+02
smoothness_worst	compactness_worst	concavity_worst

```

1.323686e-01      2.542650e-01      2.721885e-01
concave.points_worst
1.146062e-01      symmetry_worst fractal_dimension_worst

```

```
apply(wisc.data, 2, sd)
```

radius_mean	texture_mean	perimeter_mean
3.524049e+00	4.301036e+00	2.429898e+01
area_mean	smoothness_mean	compactness_mean
3.519141e+02	1.406413e-02	5.281276e-02
concavity_mean	concave.points_mean	symmetry_mean
7.971981e-02	3.880284e-02	2.741428e-02
fractal_dimension_mean	radius_se	texture_se
7.060363e-03	2.773127e-01	5.516484e-01
perimeter_se	area_se	smoothness_se
2.021855e+00	4.549101e+01	3.002518e-03
compactness_se	concavity_se	concave.points_se
1.790818e-02	3.018606e-02	6.170285e-03
symmetry_se	fractal_dimension_se	radius_worst
8.266372e-03	2.646071e-03	4.833242e+00
texture_worst	perimeter_worst	area_worst
6.146258e+00	3.360254e+01	5.693570e+02
smoothness_worst	compactness_worst	concavity_worst
2.283243e-02	1.573365e-01	2.086243e-01
concave.points_worst	symmetry_worst	fractal_dimension_worst
6.573234e-02	6.186747e-02	1.806127e-02

```
wisc.pr <- prcomp(wisc.data, scale=TRUE)
summary(wisc.pr)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	3.6444	2.3857	1.67867	1.40735	1.28403	1.09880	0.82172
Proportion of Variance	0.4427	0.1897	0.09393	0.06602	0.05496	0.04025	0.02251
Cumulative Proportion	0.4427	0.6324	0.72636	0.79239	0.84734	0.88759	0.91010
	PC8	PC9	PC10	PC11	PC12	PC13	PC14
Standard deviation	0.69037	0.6457	0.59219	0.5421	0.51104	0.49128	0.39624
Proportion of Variance	0.01589	0.0139	0.01169	0.0098	0.00871	0.00805	0.00523
Cumulative Proportion	0.92598	0.9399	0.95157	0.9614	0.97007	0.97812	0.98335
	PC15	PC16	PC17	PC18	PC19	PC20	PC21
Standard deviation	0.30681	0.28260	0.24372	0.22939	0.22244	0.17652	0.1731

```

Proportion of Variance 0.00314 0.00266 0.00198 0.00175 0.00165 0.00104 0.0010
Cumulative Proportion 0.98649 0.98915 0.99113 0.99288 0.99453 0.99557 0.9966
PC22 PC23 PC24 PC25 PC26 PC27 PC28
Standard deviation 0.16565 0.15602 0.1344 0.12442 0.09043 0.08307 0.03987
Proportion of Variance 0.00091 0.00081 0.0006 0.00052 0.00027 0.00023 0.00005
Cumulative Proportion 0.99749 0.99830 0.9989 0.99942 0.99969 0.99992 0.99997
PC29 PC30
Standard deviation 0.02736 0.01153
Proportion of Variance 0.00002 0.00000
Cumulative Proportion 1.00000 1.00000

```

Q4. From your results, what proportion of the original variance is captured by the first principal components (PC1)?

PC1 captures 0.4427 of the variance (44.27%).

Q5. How many principal components (PCs) are required to describe at least 70% of the original variance in the data?

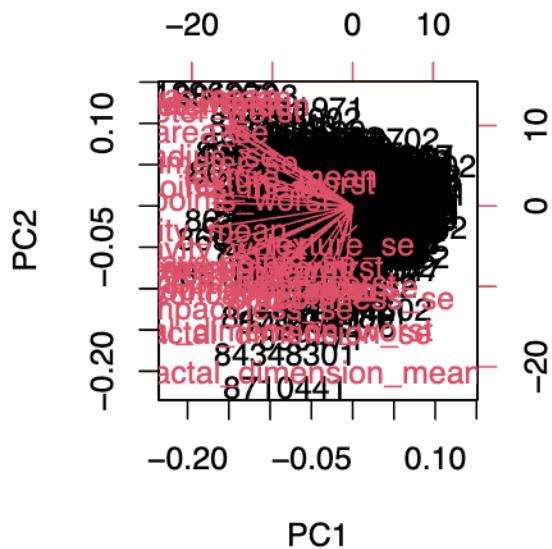
Three PCAs are required to get to at least 70% of the variance.

Q6. How many principal components (PCs) are required to describe at least 90% of the original variance in the data?

Seven PCAs are required to get to at least 90% of variance.

Let's make our main result figure-the "PC plot", or "score plot", or "ordination plot"

```
biplot(wisc.pr)
```

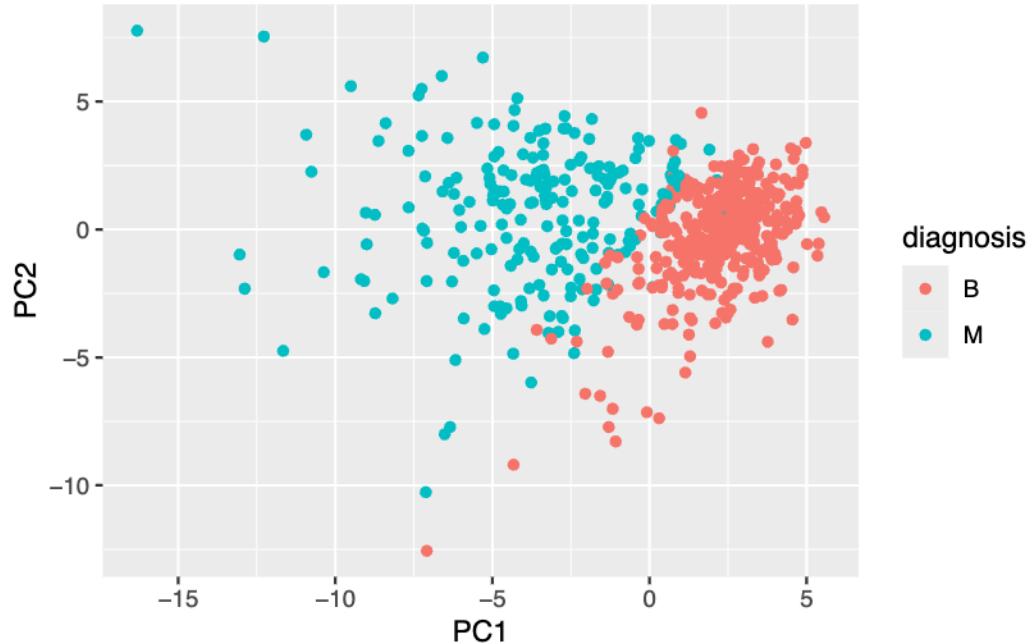


Q7. What stands out to you about this plot? Is it easy or difficult to understand? Why?

This plot is very messy and impossible to read since there is so much overlap

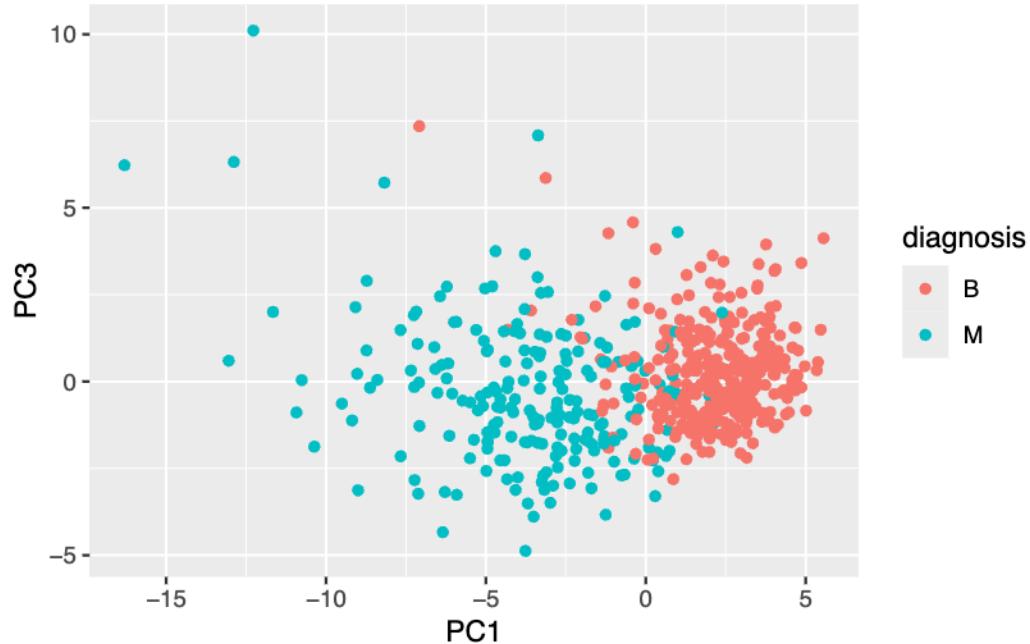
```
library(ggplot2)

ggplot(wisc.pr$x) +
  aes(PC1, PC2, col=diagnosis) +
  geom_point()
```



Q8. Generate a similar plot for principal components 1 and 3. What do you notice about these plots?

```
ggplot(wisc.pr$x)+  
  aes(PC1, PC3, col=diagnosis) +  
  geom_point()
```

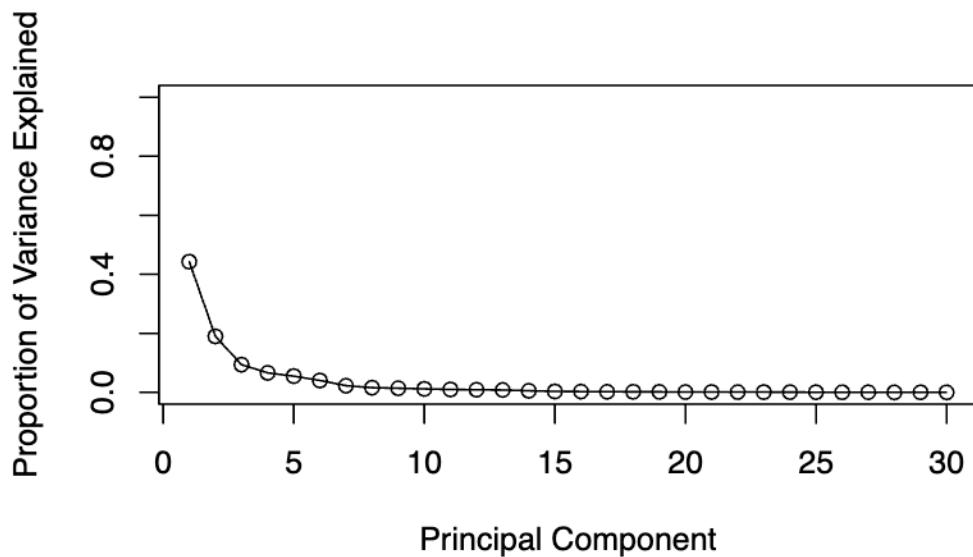


```
pr.var <- wisc.pr$sdev^2
head(pr.var)
```

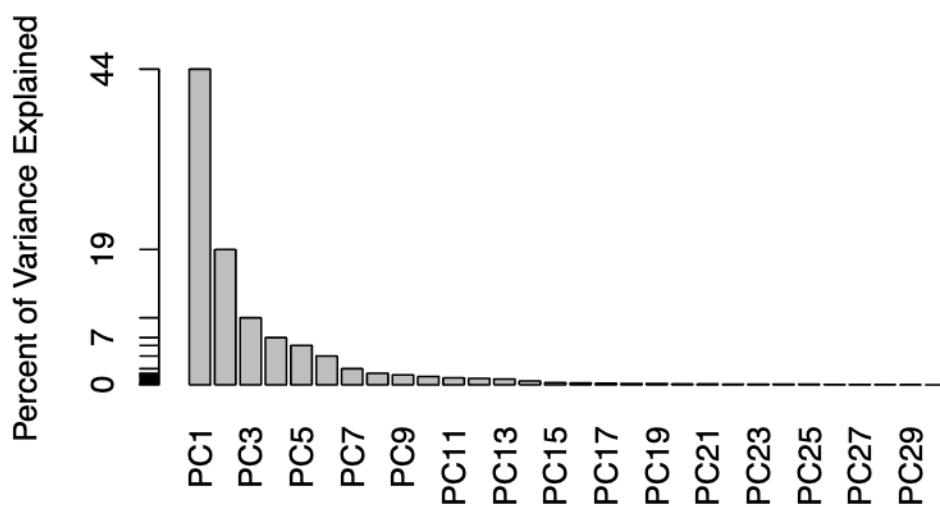
```
[1] 13.281608 5.691355 2.817949 1.980640 1.648731 1.207357
```

```
pve <- pr.var/ sum(pr.var)

plot(pve, xlab = "Principal Component",
      ylab = "Proportion of Variance Explained",
      ylim = c(0, 1), type = "o")
```



```
barplot(pve, ylab = "Percent of Variance Explained",
        names.arg=paste0("PC",1:length(pve)), las=2, axes = FALSE)
axis(2, at=pve, labels=round(pve,2)*100 )
```

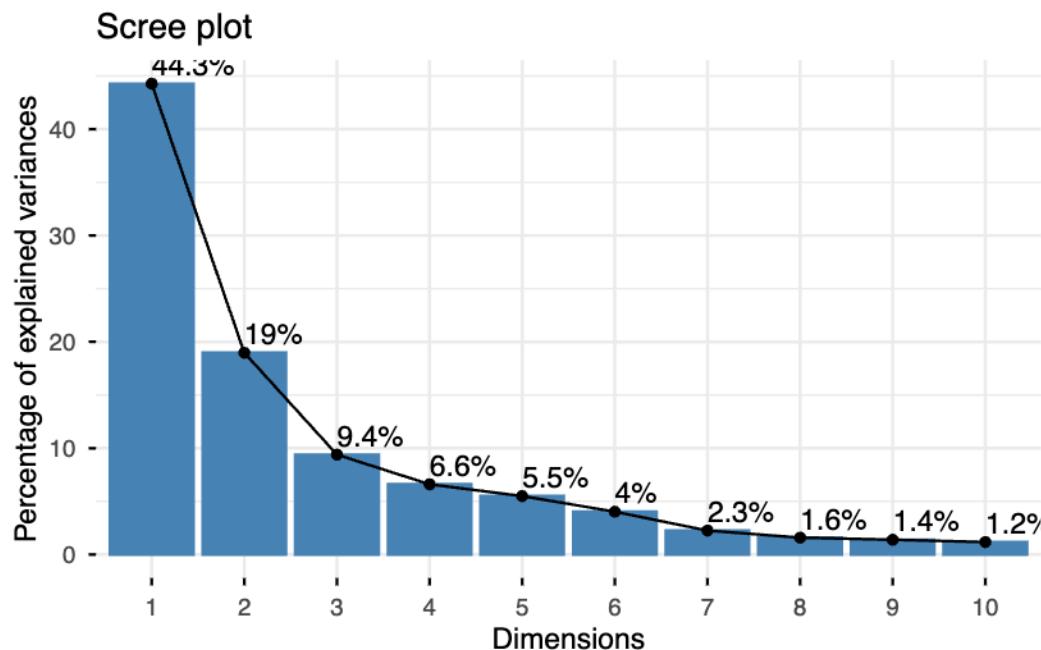


```
library(factoextra)
```

Welcome! Want to learn more? See two factoextra-related books at <https://goo.gl/ve3WBa>

```
fviz_eig(wisc.pr, addlabels = TRUE)
```

Warning in geom_bar(stat = "identity", fill = barfill, color = barcolor, :
Ignoring empty aesthetic: `width`.



Q9. For the first principal component, what is the component of the loading vector (i.e. wisc.pr\$rotation[,1]) for the feature concave.points_mean? This tells us how much this original feature contributes to the first PC.

```
wisc.pr$rotation[,1]
```

radius_mean	texture_mean	perimeter_mean
-0.21890244	-0.10372458	-0.22753729
area_mean	smoothness_mean	compactness_mean
-0.22099499	-0.14258969	-0.23928535
concavity_mean	concave.points_mean	symmetry_mean

	-0.25840048	-0.26085376	-0.13816696
fractal_dimension_mean		radius_se	texture_se
	-0.06436335	-0.20597878	-0.01742803
perimeter_se		area_se	smoothness_se
	-0.21132592	-0.20286964	-0.01453145
compactness_se		concavity_se	concave.points_se
	-0.17039345	-0.15358979	-0.18341740
symmetry_se	fractal_dimension_se		radius_worst
	-0.04249842	-0.10256832	-0.22799663
texture_worst	perimeter_worst		area_worst
	-0.10446933	-0.23663968	-0.22487053
smoothness_worst	compactness_worst		concavity_worst
	-0.12795256	-0.21009588	-0.22876753
concave.points_worst	symmetry_worst	fractal_dimension_worst	
	-0.25088597	-0.12290456	-0.13178394

-0.26085376

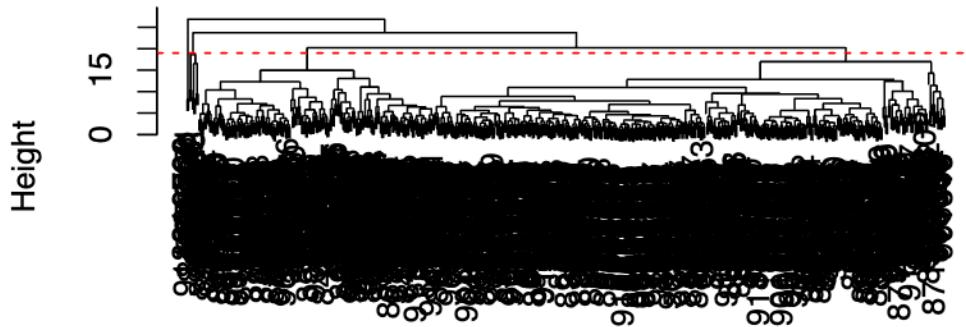
Hierarchical Clustering

```
data.scaled <- scale(wisc.data)
data.dist <- dist(data.scaled)
wisc.hclust <- hclust(data.dist, method = "complete")
```

Q10. Using the plot() and abline() functions, what is the height at which the clustering model has 4 clusters?

```
plot(wisc.hclust)
abline(h=19, col="red", lty=2)
```

Cluster Dendrogram



```
data.dist  
hclust (*, "complete")
```

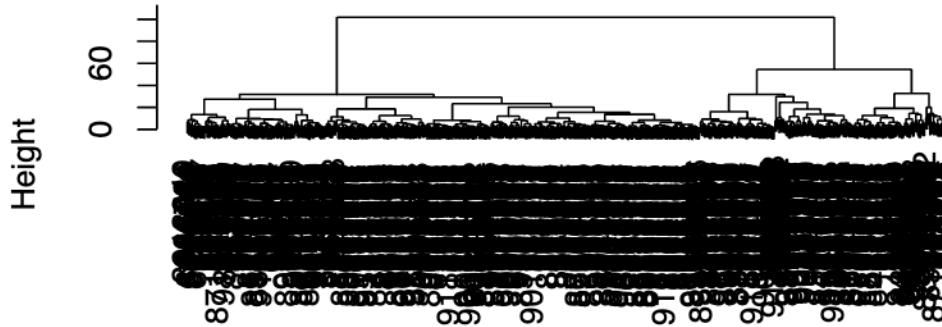
```
wisc.hclust.clusters <- cutree(wisc.hclust, h=19)  
table(wisc.hclust.clusters, diagnosis)
```

diagnosis		
wisc.hclust.clusters	B	M
1	12	165
2	2	5
3	343	40
4	0	2

Q12. Which method gives your favorite results for the same data.dist dataset?
Explain your reasoning.

```
wisc.hclust <- hclust(data.dist, method = "ward.D2")  
plot(wisc.hclust)
```

Cluster Dendrogram



```
data.dist  
hclust (*, "ward.D2")
```

I preferred the “ward.D2” method since it was the easiest method to visualize the number of clusters. The other methods ended up being a lot more chaotic (especially “single”, yuck!) making it harder to distinguish the branches.

Combining methods

```
wisc.pr.hclust <- hclust(dist(wisc.pr$x[,1:7]), method = "ward.D2")
```

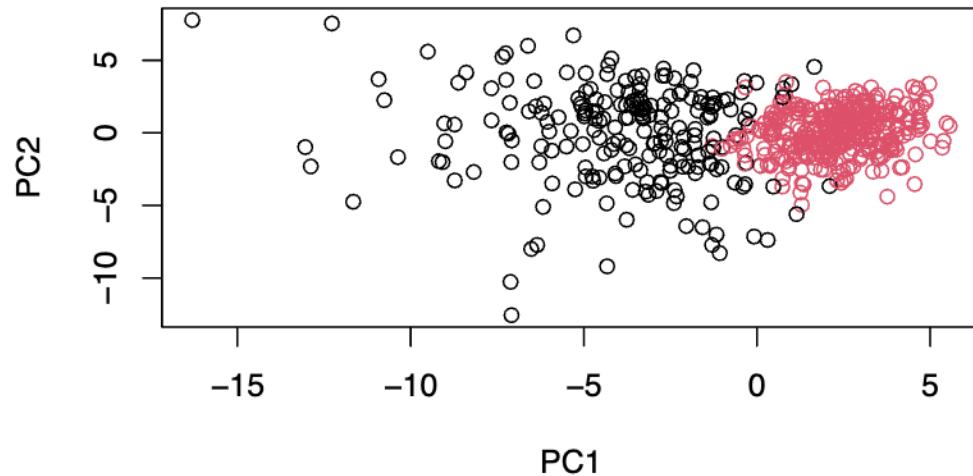
```
grps <- cutree(wisc.pr.hclust, k=2)  
table(grps)
```

```
grps  
1 2  
216 353
```

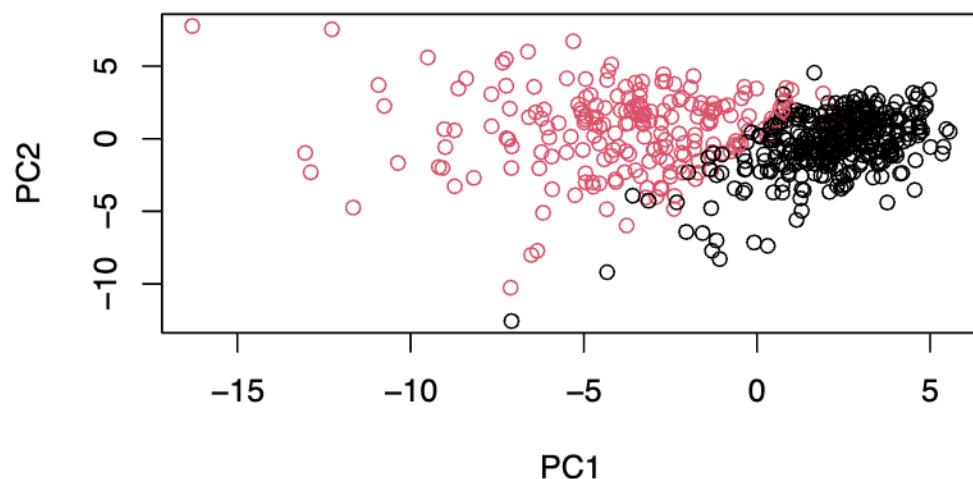
```
table(grps, diagnosis)
```

grps	B	M
1	28	188
2	329	24

```
plot(wisc.pr$x[,1:2], col=grps)
```



```
plot(wisc.pr$x[,1:2], col=diagnosis)
```



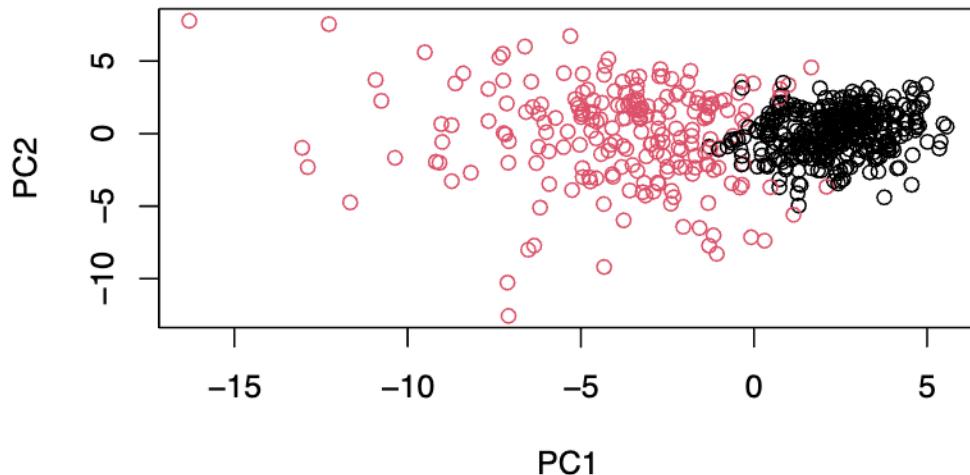
```
g <- as.factor(grps)
levels(g)
```

```
[1] "1" "2"
```

```
g <- relevel(g,2)
levels(g)
```

```
[1] "2" "1"
```

```
plot(wisc.pr$x[,1:2], col=g)
```



```
wisc.pr.hclust <- hclust(dist(wisc.pr$x[, 1:7]), method="ward.D2")
```

```
wisc.pr.hclust.clusters <- cutree(wisc.pr.hclust, k=2)
```

```
table(wisc.pr.hclust.clusters)
```

```
wisc.pr.hclust.clusters
 1   2
216 353
```

Q13. How well does the newly created model with four clusters separate out the two diagnoses?

```
table(wisc.pr.hclust.clusters, diagnosis)
```

wisc.pr.hclust.clusters	B	M	diagnosis
1	28	188	
2	329	24	

It separates them out pretty well, but not perfectly. It captured 88.7% of the malignant samples, which in a clinical practice seems really low. It captured 92% of the benign tumors which is a little better.

Q14. How well do the hierarchical clustering models you created in previous sections (i.e. before PCA) do in terms of separating the diagnoses? Again, use the table() function to compare the output of each model (wisc.km\$cluster and wisc.hclust.clusters) with the vector containing the actual diagnoses.

```
table(wisc.hclust.clusters, diagnosis)
```

wisc.hclust.clusters	B	M	diagnosis
1	12	165	
2	2	5	
3	343	40	
4	0	2	

It does a little worse than the previous method in identifying the malignant tumors, but correctly identified more of the benign tumors into one group. So it is identifying more malignant tumors as benign which is more dangerous in practice since people could go untreated if this algorithm was determining their sample results.

Prediction

```
url <- "https://tinyurl.com/new-samples-CSV"  
new <- read.csv(url)  
npc <- predict(wisc.pr, newdata=new)  
npc
```

```

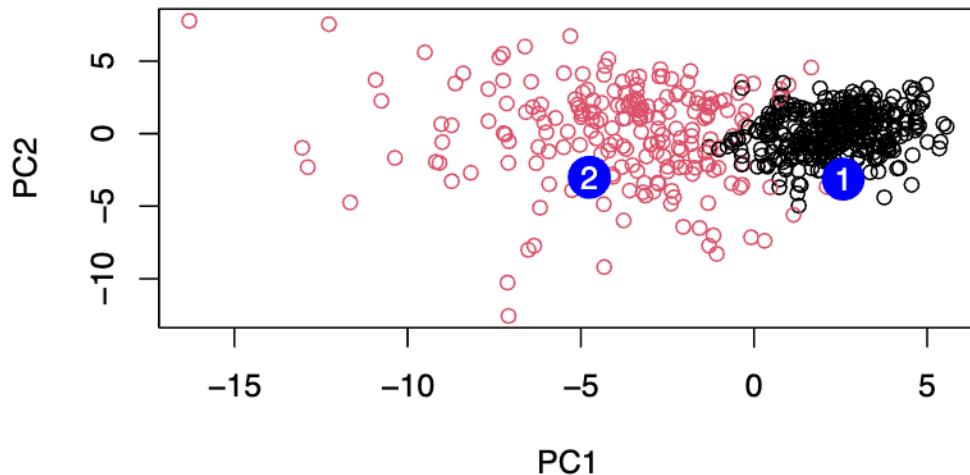
          PC1      PC2      PC3      PC4      PC5      PC6      PC7
[1,]  2.576616 -3.135913 1.3990492 -0.7631950 2.781648 -0.8150185 -0.3959098
[2,] -4.754928 -3.009033 -0.1660946 -0.6052952 -1.140698 -1.2189945  0.8193031
          PC8      PC9      PC10     PC11     PC12     PC13     PC14
[1,] -0.2307350 0.1029569 -0.9272861 0.3411457 0.375921 0.1610764 1.187882
[2,] -0.3307423 0.5281896 -0.4855301 0.7173233 -1.185917 0.5893856 0.303029
          PC15     PC16     PC17     PC18     PC19     PC20
[1,]  0.3216974 -0.1743616 -0.07875393 -0.11207028 -0.08802955 -0.2495216
[2,]  0.1299153  0.1448061 -0.40509706  0.06565549  0.25591230 -0.4289500
          PC21     PC22     PC23     PC24     PC25     PC26
[1,]  0.1228233  0.09358453  0.08347651  0.1223396  0.02124121  0.078884581
[2,] -0.1224776  0.01732146  0.06316631 -0.2338618 -0.20755948 -0.009833238
          PC27     PC28     PC29     PC30
[1,]  0.220199544 -0.02946023 -0.015620933  0.005269029
[2,] -0.001134152  0.09638361  0.002795349 -0.019015820

```

```

plot(wisc.pr$x[,1:2], col=g)
points(npc[,1], npc[,2], col="blue", pch=16, cex=3)
text(npc[,1], npc[,2], c(1,2), col="white")

```



Q16. Which of these new patients should we prioritize for follow up based on your results?

We should prioritize patients from group 2 for followup since they have been clustered as the malignant group and therefore have sample results that more closely relate to the malignant features.

```
sessionInfo()
```

```
R version 4.5.1 (2025-06-13)
Platform: aarch64-apple-darwin20
Running under: macOS Sequoia 15.6.1

Matrix products: default
BLAS:      /Library/Frameworks/R.framework/Versions/4.5-arm64/Resources/lib/libRblas.0.dylib
LAPACK:    /Library/Frameworks/R.framework/Versions/4.5-arm64/Resources/lib/libRlapack.dylib; 

locale:
[1] en_US.UTF-8/en_US.UTF-8/en_US.UTF-8/C/en_US.UTF-8/en_US.UTF-8

time zone: America/Los_Angeles
tzcode source: internal

attached base packages:
[1] stats      graphics   grDevices utils      datasets   methods    base

other attached packages:
[1] factoextra_1.0.7 ggplot2_4.0.0

loaded via a namespace (and not attached):
 [1] gtable_0.3.6       jsonlite_2.0.0      dplyr_1.1.4        compiler_4.5.1
 [5] ggsignif_0.6.4     tidyselect_1.2.1    Rcpp_1.1.0         tidyrr_1.3.1
 [9] scales_1.4.0       yaml_2.3.10       fastmap_1.2.0     R6_2.6.1
[13] ggpubr_0.6.2       labeling_0.4.3     generics_0.1.4    Formula_1.2-5
[17] knitr_1.50          backports_1.5.0    ggrepel_0.9.6     tibble_3.3.0
[21] car_3.1-3           pillar_1.11.1     RColorBrewer_1.1-3 rlang_1.1.6
[25] broom_1.0.10        xfun_0.53        S7_0.2.0          cli_3.6.5
[29] withr_3.0.2          magrittr_2.0.4    digest_0.6.37     grid_4.5.1
[33] lifecycle_1.0.4      vctrs_0.6.5       rstatix_0.7.3    evaluate_1.0.5
[37] glue_1.8.0            farver_2.1.2     abind_1.4-8       carData_3.0-5
[41] rmarkdown_2.30        purrr_1.1.0       tools_4.5.1       pkgconfig_2.0.3
[45] htmltools_0.5.8.1
```